

Metasynth: Transparent Generation of Synthetic Tabular Data with Privacy Guarantees

Raoul Schram^{1*}, Thom Volker^{1,3}, Ayoub Bagheri¹, Ron Scholten¹, Samuel Spithorst¹, and Erik-Jan van Kesteren^{1,2*}

¹ Utrecht University, The Netherlands ² ODISSEI: Open Data Infrastructure for Social Science and Economic Innovations, The Netherlands ³ Statistics Netherlands ¶ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ¶
- [Repository](#) ¶
- [Archive](#) ¶

Editor: [Open Journals](#) ¶

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Synthetic data is a promising tool for improving the accessibility of datasets that are otherwise too sensitive to be shared publicly. To this end, we introduce metasynth, a Python package for generating synthetic data from tabular datasets. Unlike existing synthetic data generation software, metasynth is built on a simple generative model with a “naive” marginal independence assumption — an explicit choice that lowers the multivariate precision of the synthetic data in order to maintain transparency and auditability, to keep information leakage to a minimum, and even to enable privacy or disclosure risk guarantees through a plug-in system. While the analytical validity of the generated data is thus intentionally limited, its potential uses are broad, including exploratory analyses, code development and testing, and external communication and teaching. Metasynth is flexible, scalable, and easily extended to meet diverse privacy needs.



Figure 1: Logo of the metasynth project.

Statement of need

Metasynth is a python package for generating synthetic data with a focus on privacy and disclosure control. It is aimed at owners of sensitive datasets such as public organisations, research groups, and individual researchers who want to improve the accessibility of their data for research and reproducibility by others. The goal of metasynth is to make it easy for data owners to share the structure and and approximation of the content of their data with others without any privacy concerns.

With this goal in mind, metasynth distinguishes itself from existing software for generating synthetic data (e.g., [Nowok et al., 2016](#); [Ping et al., 2017](#); [Templ et al., 2017](#)) by restricting itself to the “augmented plausible” category of synthetic data ([Bates et al., 2019](#)). This choice enables the software to generate synthetic data with **privacy and disclosure guarantees** through a plug-in system. Moreover, our system provides an **auditable and editable intermediate representation** in the form of a human- and machine-readable .json metadata file from which new data can be synthesized.

Through our focus on privacy and transparency, metasynt explicitly avoids generating synthetic data with high analytical validity. The data generated by our system is realistic in terms of data structure and plausible in terms of values for each variable, but any multivariate relations or conditional patterns are excluded. This has implications for how this synthetic data can be used: not for statistical analysis and inference, but rather for initial exploration, analysis script development, and communication outside the data owner's institution. In the intended use case, an external researcher can make use of the synthetic data to assess the feasibility of their intended research before making the (often time-consuming) step of requesting access to the sensitive source data for the final analysis.

As mentioned before, the privacy capacities of metasynt are extensible through a plug-in system, recognizing that different data owners have different needs and definitions of privacy. A data owner can define under which conditions they would accept open distribution of their synthetic data — be it based on differential privacy (Dwork, 2006), statistical disclosure control (Wolf, 2012), k-anonymity (Sweeney, 2002), or another specific definition of privacy. As part of the initial release of metasynt, we publish two proof-of-concept plugins: one following the disclosure control guidelines from Eurostat (Bond et al., 2015), and one based on the sample-and-aggregate technique for differential privacy (Dwork & Smith, 2010, p. 142).

Software features

At its core, Metasynt is designed for three functions, which are briefly described in this section:

1. **Estimation:** Automatically select univariate distributions and fit them to a well-defined tabular dataset, possibly with privacy guarantees.
2. **(De)serialization:** Create an intermediate representation of the fitted model for auditing, editing, and exporting.
3. **Generation:** Generate new synthetic datasets based on the fitted model or its serialized representation.

Estimation

The generative model for multivariate datasets in metasynt makes the simplifying assumption of marginal independence: each column is considered separately, just as is done in e.g., naïve Bayes classifiers (Hastie et al., 2009). Formally, this leads to the following generative model for the K -variate data \mathbf{x} :

$$p(\mathbf{x}) = \prod_{k=1}^K p(x_k) \quad (1)$$

There are many advantages to this naïve approach when compared to more advanced generative models: it is transparent and explainable, it is able to flexibly handle data of mixed types, and it is computationally scalable to high-dimensional datasets. As mentioned before, the tradeoff is the limited analytical validity when the independence assumption does not hold: in the synthetic data, the expected value of correlations, regression parameters, and other measures of association is 0.

Model estimation starts with an appropriately pre-processed data frame. For metasynt, this means the data frame is tidy (Wickham, 2014), each column has the correct data type, and missing data are represented by a missing value. Internally, our software uses the polars data frame library (Vink et al., 2023), as it is performant, has consistent data types, and native support for missing data (null). A simple example source table could look like this (note that categorical data has the appropriate cat data type, not str):

ID	fruits	B	cars	optional

---	---	---	---	---
i64	cat	i64	cat	i64
1	banana	5	beetle	28
2	banana	4	audi	300
3	apple	3	beetle	null
4	apple	2	beetle	2
5	banana	1	beetle	-30

For each data type supported by metasynt, there is a set of candidate distributions that can be fitted to that data type (see Table Table 1). To estimate the generative model of Equation Equation 1, for each variable the software fits all compatible candidate distributions — by default with maximum likelihood estimation — and then selects the one with the lowest AIC (Akaike, 1973).

Table 1: Candidate distributions associated with data types in the core metasynt package.

Variable type	Data type	example	candidate distributions
continuous	float	1.0, 2.1, ...	UniformDistribution, NormalDistribution, ...
discrete	int	1, 2, ...	DiscreteUniformDistribution
categorical	pl.Categorical	gender, country	MultinoulliDistribution
structured string	str	Room number A108, C122	RegexDistribution
unstructured string	str	Names, open answers	FakerDistribution, LLMDistribution
temporal	Date, Datetime	2021-01-13, 01:40:12	DateUniformDistribution

Generative model estimation done in code as follows:

```
from metasynt import MetaDataset

# "ID" column is the primary key;
# ensure that it has unique values
# when data is synthesized later
spec_dict = {
    "ID": {"unique": True}
}

# create metadataset
mds = MetaDataset.from_dataframe(df, spec=spec_dict)
```

Serialization and deserialization

After a fitted model object is created, metasynt allows it to be transparently stored in a human- and machine-readable .json file. This file can be considered as metadata: it contains dataset-level descriptive information as well as variable-level information. The metadata format has a specific structure, which we call generative metadata format, or gmf. The header contains the following dataset-level information:

```

"n_rows": 5,
"n_columns": 5,
"provenance": {
  "created by": {
    "name": "MetaSynth",
    "version": "0.1.2+15.ged3af36",
    "privacy": null
  },
  "creation time": "2022-11-17T13:54:16.686166"
}

```

98 Then, for each variable the gmf file contains information the name, the data type, the
 99 proportion of missing values, and the distribution fitted on the data. For example, a table
 100 column containing different types of fruits could result in the following .json:

```

{
  "name": "fruits",
  "type": "categorical",
  "dtype": "<class 'polars.datatypes.Categorical'>",
  "prop_missing": 0.0,
  "distribution": {
    "name": "MultinoulliDistribution",
    "parameters": {
      "labels": ["apple", "banana"],
      "probs": [0.4, 0.6]
    }
  }
}

```

101 There are several advantages to creating such a serialized representation. First, it can be
 102 audited: the data owner can see exactly what information from the real data is made public
 103 through exporting the synthetic data, namely, the parameters of the distribution. Second,
 104 the file can be edited. For example, if a data owner thinks some of the labels of the “fruit”
 105 column contain sensitive information, these can simply be pseudonymized in the metadata
 106 file. Third, after exporting this file, an unlimited number of synthetic records can be created
 107 without incurring additional privacy risks, because the original data is no longer part of the
 108 synthetization process.

109 Data generation

110 After creating either the fitted model object from the original data or by deserializing a model
 111 object from a .json file, new data can be generated by the object. For each variable in the
 112 model object, the software randomly samples from the fitted distribution to create a synthetic
 113 version of the data.

```

from metasynt import MetaDataset

# load json into a metadataset object
mds = MetaDataset.from_json("metasynt_example.json")

# create a fake dataset
mds.synthesize(10)

```

114 This may result in the following polars data frame. Note that missing values in the optional
 115 column are appropriately reproduced as well, courtesy of the “prop_missing” entry in the
 116 metadata format.

117 shape: (10, 5)

118					
119	ID	fruits	B	cars	optional
120	---	---	---	---	---
121	i64	cat	i64	cat	i64
122					
123	1	banana	4	beetle	null
124	2	banana	3	audi	null
125	3	banana	1	beetle	223
126	4	banana	0	beetle	258
127
128	7	banana	3	beetle	298
129	8	banana	2	beetle	67
130	9	banana	4	beetle	-30
131	10	banana	2	beetle	172
132					

Plug-ins and automatic privacy

In addition to the core features described above, the `metasynth` package allows for plug-ins: add-on packages that alter the behaviour of the parameter estimation. Through this system, privacy guarantees can be built into `metasynth`. For example, a package called `metasynth-disclosure-control` implements the disclosure control output guidelines from Eurostat (Bond et al., 2015) by re-implementing the `fit()` method of the candidate distributions shown in Table Table 1 to include a micro-aggregation step. In this way, information transfer from the sensitive real data to the synthetic public data can be further reduced.

This plug-in system is user-friendly: the user only needs to `pip install` the package and then `metasynth` can automatically find it to make the methods accessible:

```
from metasynth import MetaDataset
from metasynthcontrib.disclosure import DisclosurePrivacy

mds = MetaDataset.from_dataframe(df, privacy=DisclosurePrivacy())
```

Conclusion

Synthetic data is a valuable tool for communicating about sensitive datasets. In this work, we have presented the software `metasynth`, which allows data owners to generate a synthetic version of their sensitive tabular data with a focus on privacy and transparency. Unlike existing tools for generating synthetic data, we choose to aim for low analytic validity to enable high privacy guarantees: the underlying model makes a simplifying independence assumption, resulting in few parameters and thus a very small information transfer. This approach additionally allows for disclosure guarantees through a plug-in system.

Further documentation and examples can be found on metasynth.readthedocs.io.

Acknowledgements

This research was conducted in whole or in part using ODISSEI, the Open Data Infrastructure for Social Science and Economic Innovations (<https://ror.org/03m8v6t10>)

The `metasynth` project is supported by the FAIR Research IT Innovation Fund of Utrecht University (as of March 2023)

References

- Akaike, H. (1973). Information theory and an extension of the maximum likelihood principle. *2nd International Symposium on Information Theory*, 267–281.
- Bates, A., Spakulová, I., Dove, I., & Meador, A. (2019). *ONS methodology working paper series number 16—synthetic data pilot*.
- Bond, S., Brandt, M., & Wolf, P. de. (2015). *Guidelines for output checking*. eurostat.
- Dwork, C. (2006). Differential privacy. *International Colloquium on Automata, Languages, and Programming*, 1–12.
- Dwork, C., & Smith, A. (2010). Differential privacy for statistics: What we know and what we want to learn. *Journal of Privacy and Confidentiality*, 1(2).
- Hastie, T., Tibshirani, R., Friedman, J. H., & Friedman, J. H. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (Vol. 2). Springer.
- Nowok, B., Raab, G. M., & Dibben, C. (2016). Synthpop: Bespoke creation of synthetic data in R. *Journal of Statistical Software*, 74, 1–26.
- Ping, H., Stoyanovich, J., & Howe, B. (2017). Datasynthesizer: Privacy-preserving synthetic datasets. *Proceedings of the 29th International Conference on Scientific and Statistical Database Management*, 1–5.
- Sweeney, L. (2002). K-anonymity: A model for protecting privacy. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, 10(05), 557–570.
- Templ, M., Meindl, B., Kowarik, A., & Dupriez, O. (2017). Simulation of synthetic complex data: The R package simPop. *Journal of Statistical Software*, 79(10), 1–38.
- Vink, R., Gooijer, S. de, Beedie, A., Gorelli, M. E., Zundert, J. van, Hulselmans, G., Grinstead, C., Santamaria, M., Heres, D., ibENPC, Magarick, J., Leitao, J., Marshall, Wilksch, M., Heerden, M. van, Borchert, O., Jermain, C., Peek, J., Russell, R., ... Robert. (2023). *Pola-rs/polars: Python polars 0.18.8 (py-0.18.8)*. Zenodo. <https://doi.org/10.5281/zenodo.8167449>
- Wickham, H. (2014). Tidy data. *Journal of Statistical Software*, 59(10), 1–23. <https://doi.org/10.18637/jss.v059.i10>
- Wolf, P.-P. de. (2012). *Statistical disclosure control*. Wiley & Sons, Chichester.